

Dynamic Indexes


Algorithmic Design

Alberto Casagrande

Email: `acasagrande@units.it`

a.y. 2020/2021

Motivations



A Simple Problem for Registry Office

Let us consider the registry office

For each newborn, they record a set of data e.g., name, birthday, parents, etc.

So, the registry data-set (hopefully) changes quite often

What if they frequently perform a birthday-based search on the data-set? E.g., Find all the baby born a given (variable) day?

Some Possible Strategies

They may:

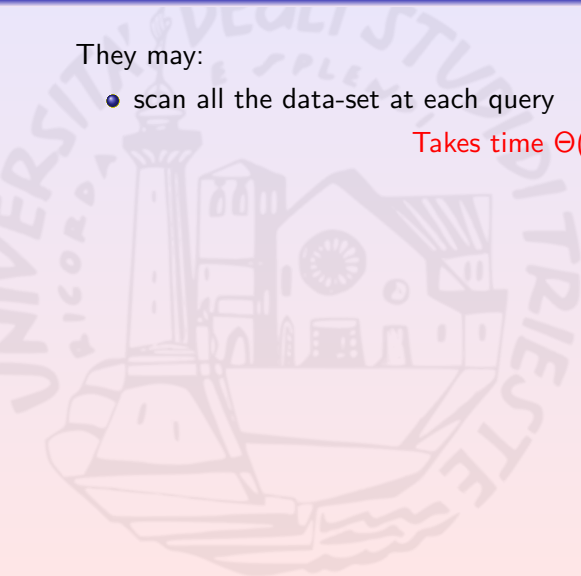


Some Possible Strategies

They may:

- scan all the data-set at each query

Takes time $\Theta(n)$



They may:

- scan all the data-set at each query

Takes time $\Theta(n)$

- sort the data-set by birthday at each insertion

Takes time $\Theta(n)$ per insertion (Radix Sort) + $\Theta(\log n)$ per query. What if they also want to perform searches by name?

They may:

- scan all the data-set at each query
Takes time $\Theta(n)$
- sort the data-set by birthday at each insertion
Takes time $\Theta(n)$ per insertion (Radix Sort) + $\Theta(\log n)$ per query. What if they also want to perform searches by name?
- sort the data-set by birthday at each query
Takes time $\Theta(n + \log n) = \Theta(n)$ per query
(Radix Sort and DS)

They may:

- scan all the data-set at each query

Takes time $\Theta(n)$

- sort the data-set by birthday at each insertion

Takes time $\Theta(n)$ per insertion (Radix Sort) + $\Theta(\log n)$ per query. What if they also want to perform searches by name?

- sort the data-set by birthday at each query

Takes time $\Theta(n + \log n) = \Theta(n)$ per query
(Radix Sort and DS)

- They may:
 - scan all the data-set at each query
Takes time $\Theta(n)$
 - sort the data-set by birthday at each insertion
Takes time $\Theta(n)$ per insertion (Radix Sort) + $\Theta(\log n)$ per query. What if they also want to perform searches by name?
 - sort the data-set by birthday at each query
Takes time $\Theta(n + \log n) = \Theta(n)$ per query
(Radix Sort and DS)

The data-set often changes, thus, array is not the most suitable data structure to achieve this goal

They may:

- scan all the data-set at each query

Takes time $\Theta(n)$

- sort the data-set by birthday at each insertion

Takes time $\Theta(n)$ per insertion (Radix Sort) + $\Theta(\log n)$ per query. What if they also want to perform searches by name?

- sort the data-set by birthday at each query

Takes time $\Theta(n + \log n) = \Theta(n)$ per query
(Radix Sort and DS)

The data-set often changes, thus, array is not the most suitable data structure to achieve this goal

A Dynamic Data Structure for Indexing

We need a data structure providing (efficient) support for:

- adding new data
- searching data
- removing data

A Dynamic Data Structure for Indexing

We need a data structure providing (efficient) support for:

- adding new data
- searching data
- removing data

We are aiming to build an **index** i.e., an auxiliary data structure to “efficiently” perform above operations

Binary Search Trees

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

`x.parent` is the parent of `x`

Binary Search Tree Notation

We have already introduced trees as ADT.

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

`x.parent` is the parent of `x`

`x.key` is the value stored in `x` i.e., its key

Binary Search Tree Notation

We have already introduced trees as ADT.

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

`x.parent` is the parent of `x`

x.key is the value stored in x i.e., its key

`x.left=NIL`, then `x` misses the left child

Binary Search Tree Notation

We have already introduced trees as ADT.

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

`x.parent` is the parent of `x`

x.key is the value stored in x i.e., its key

x.left=NIL, then x misses the left child

T.root is the root of the tree T and $T.root.parent = \text{NIL}$

```
def IS_ROOT(x):
    return x.parent=NIL
enddef

def IS_RIGHT_CHILD(x):
    return  $\neg$ IS_ROOT(x) and x.parent.right=x
enddef

def SIBLING(x):    # get x's sibling
    if IS_RIGHT_CHILD(x):
        return x.parent.left
    endif
    return x.parent.right
enddef
```

```
def CHILDHOOD_SIDE(x): # get x's side w.r.t.
                        # its parent
    if IS_RIGHT_CHILD(x):
        return RIGHT
    endif
    return LEFT
enddef

def REVERSE_SIDE(side): # reverse the side
    if side=LEFT:
        return RIGHT
    endif
    return LEFT
enddef
```

Some Useful $\Theta(1)$ Functions (Cont'd 2)

```
def GET_CHILD(x, side):  # get x's child on side
    if side==LEFT:
        return x.left
    endif

    return x.right
enddef
```

Some Useful $\Theta(1)$ Functions (Cont'd 3)

```

def SET_CHILD(x, side, y):  # set x's child
    if side=LEFT:
        x.left ← y
    else:
        x.right ← y
    endif

    if y≠NIL:
        y.parent ← x
    endif
enddef

```

Some Useful $\Theta(1)$ Functions (Cont'd 4)

```
def UNCLE(x):      #get x's uncle  
    return SIBLING(x.parent)  
enddef
```

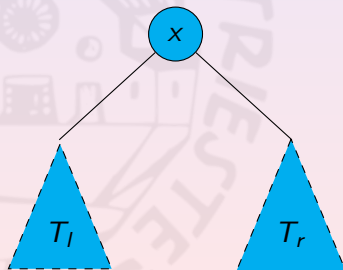
```
def GRANDPARENT(x):  # get x's granpa  
    return x.parent.parent  
enddef
```

```
def NEW_NODE(v):      # build a new node  
    x.key  $\leftarrow$  v  
    x.parent  $\leftarrow$  NIL  
    x.right  $\leftarrow$  NIL  
    x.left  $\leftarrow$  NIL  
enddef
```

Binary Search Trees

A **Binary Search Tree (BST)** is a tree s.t.:

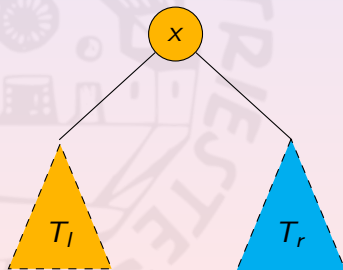
- all the keys belong to a totally ordered set w.r.t. \preceq



Binary Search Trees

A **Binary Search Tree (BST)** is a tree s.t.:

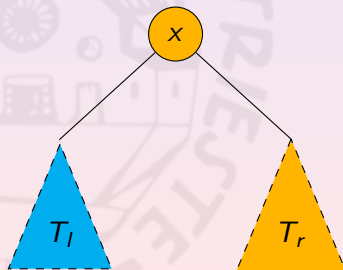
- all the keys belong to a totally ordered set w.r.t. \preceq
- if x_l is in the left sub-tree of x , then $x_l.key \preceq x.key$



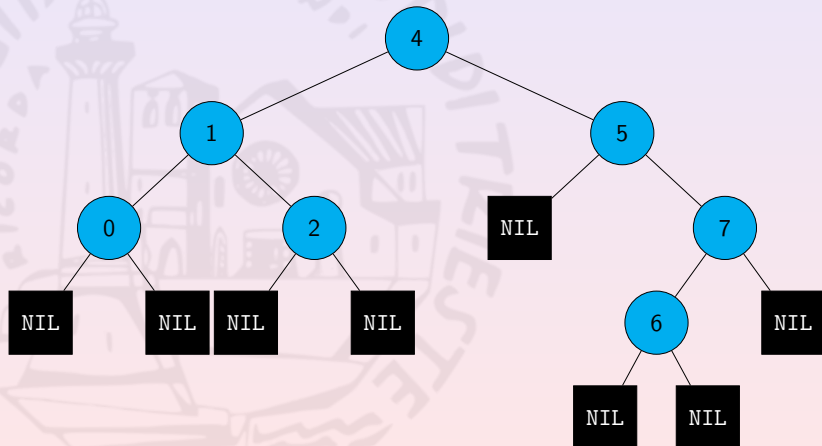
Binary Search Trees

A **Binary Search Tree (BST)** is a tree s.t.:


- all the keys belong to a totally ordered set w.r.t. \preceq
- if x_l is in the left sub-tree of x , then $x_l.key \preceq x.key$
- if x_r is in the right sub-tree of x , then $x.key \preceq x_r.key$



Binary Search Trees: an Example



Binary Search Trees: an Example




In-Order Walk

```
def INORDER_WALK_AUX(x):  
    if x ≠ NIL:  
        INORDER_WALK_AUX(x.left)  
        print x.key  
        INORDER_WALK_AUX(x.right)  
    endif  
endif  
  
def INORDER_WALK(T):  
    INORDER_WALK_AUX(T.root)  
endif
```

Searching for the Maximum/Minimum

Due to the BST property:

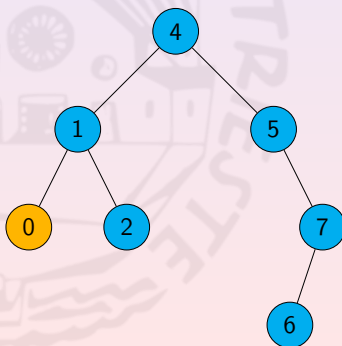
- the minimum key is contained by the first node on the leftmost branch that has not a left child



Searching for the Maximum/Minimum

Due to the BST property:

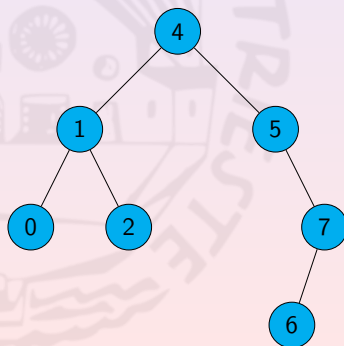
- the minimum key is contained by the first node on the leftmost branch that has not a left child



Searching for the Maximum/Minimum

Due to the BST property:

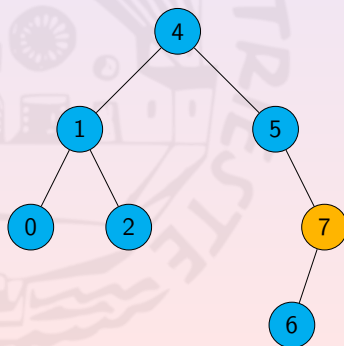
- the minimum key is contained by the first node on the leftmost branch that has not a left child
- the maximum key is contained by the first node on the rightmost branch that has not a right child



Searching for the Maximum/Minimum

Due to the BST property:

- the minimum key is contained by the first node on the leftmost branch that has not a left child
- the maximum key is contained by the first node on the rightmost branch that has not a right child



Searching for the Maximum/Minimum: Pseudo-Code

```
def MINIMUM_IN_SUBTREE(x):  
    while x.left  $\neq$  NIL:  
        x  $\leftarrow$  x.left  
    endif  
    return x  
endif
```

```
def MAXIMUM_IN_SUBTREE(x):  
    while x.right  $\neq$  NIL:  
        x  $\leftarrow$  x.right  
    endif  
    return x  
endif
```

Searching for the Maximum/Minimum: Pseudo-Code

```
def MINIMUM_IN_SUBTREE(x):  
    while x.left ≠ NIL:  
        x ← x.left  
    endif  
    return x  
endif
```

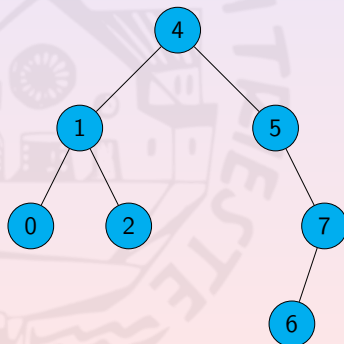
 $O(h_T)$

```
def MAXIMUM_IN_SUBTREE(x):  
    while x.right ≠ NIL:  
        x ← x.right  
    endif  
    return x  
endif
```

 $O(h_T)$

Successor of a Node

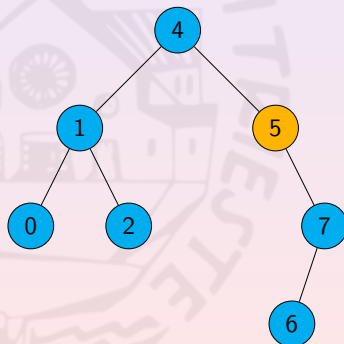
Due to the BST property, the successor w.r.t. \preceq of n is either



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either

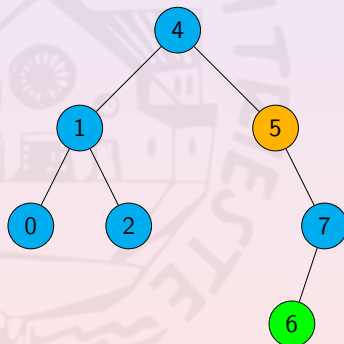
- the node containing the minimum in the right sub-tree of n



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either

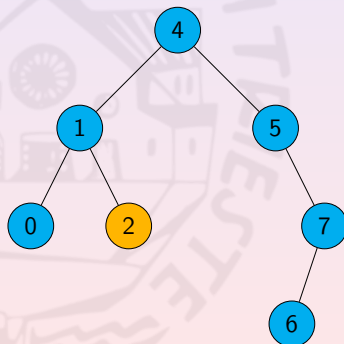
- the node containing the minimum in the right sub-tree of n or



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either

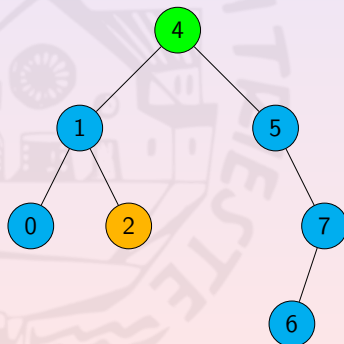
- the node containing the minimum in the right sub-tree of n or
- the nearest “right-ancestor” of n , if n has not right child



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either

- the node containing the minimum in the right sub-tree of n or
- the nearest “right-ancestor” of n , if n has not right child



Successor of a Node: Pseudo-Code

```
def SUCCESSOR(x):  
    if x.right  $\neq$  NIL:  
        return MINIMUM_IN_SUBTREE(x.right)  
    endif  
  
    y  $\leftarrow$  x.parent  
    while y  $\neq$  NIL and IS_RIGHT_CHILD(x):  
        x  $\leftarrow$  y  
        y  $\leftarrow$  x.parent  
    endwhile  
  
    return y  
enddef
```


Successor of a Node: Pseudo-Code

```
def SUCCESSOR(x):  
    if x.right ≠ NIL:  
        return MINIMUM_IN_SUBTREE(x.right)  
    endif  
  
    y ← x.parent  
    while y ≠ NIL and IS_RIGHT_CHILD(x):  
        x ← y  
        y ← x.parent  
    endwhile  
  
    return y  
enddef
```


$O(h_T)$

Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root x :

- if n is NIL or $x.key = v$, return n
- if $x.key \leq v$, search on the right sub-tree
- if $x.key \not\leq v$, search on the left sub-tree

Searching for 3

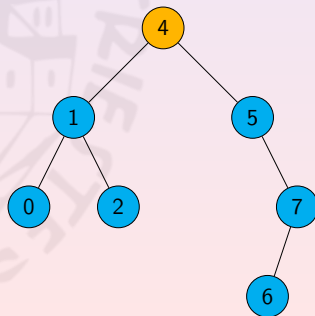


Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root x :

- if n is NIL or $x.key = v$, return n
- if $x.key \leq v$, search on the right sub-tree
- **if $x.key \not\leq v$, search on the left sub-tree**

Searching for 3

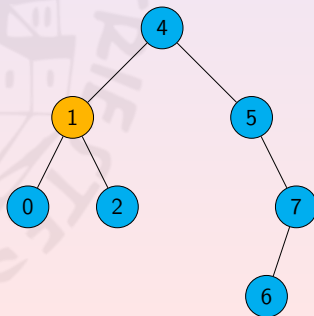


Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root x :

- if n is NIL or $x.key = v$, return n
- **if $x.key \preceq v$, search on the right sub-tree**
- if $x.key \not\preceq v$, search on the left sub-tree

Searching for 3




Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root x :

- if n is NIL or $x.key = v$, return n
- **if $x.key \preceq v$, search on the right sub-tree**
- if $x.key \not\preceq v$, search on the left sub-tree

Searching for 3




Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root x :

- if n is NIL or $x.key = v$, return n
- if $x.key \leq v$, search on the right sub-tree
- if $x.key \not\leq v$, search on the left sub-tree

Searching for 3



Searching for a Value in a BST: Pseudo-Code

```
def SEARCH_SUBTREE(x, v):  
    while x ≠ NIL:  
        if x.key ≤ v:  
            if v ≤ x.key:  
                return x  
            endif  
            x ← x.right  
        else:  
            x ← x.left  
        endif  
    endwhile  
enddef
```

Searching for a Value in a BST: Complexity

Each iteration performs $\Theta(1)$ operations

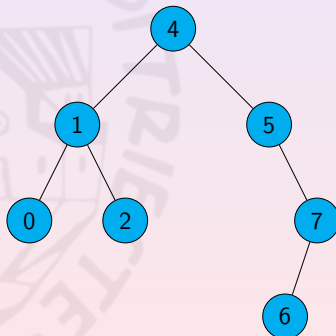
The # of iterations depends on the height h_T of T and on v

The algorithm takes time $O(h_T)$

Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key


Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key

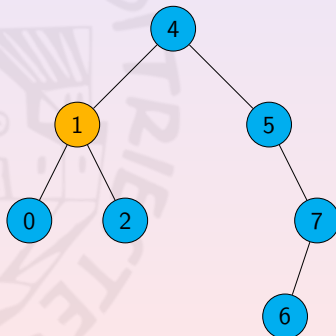
Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key

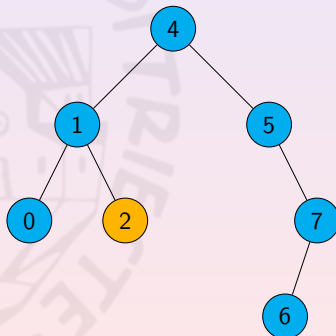
Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key


Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key

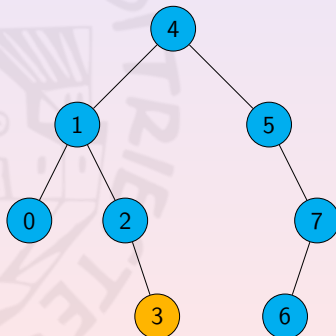
Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key

Inserting 3



Inserting a Value in a BST: Pseudo-Code

```

def INSERT_BST(T,v): # v is the new value
    x  $\leftarrow$  T.root
    y  $\leftarrow$  NIL    # y is x's parent

    # search the right place for z
    while x  $\neq$  NIL:
        y  $\leftarrow$  x
        if v  $\preceq$  x.key:
            if x.key  $\preceq$  v:
                return HANDLE.MULTI.INSERT(x,v)
            endif
            x  $\leftarrow$  x.left
        else:
            x  $\leftarrow$  x.right
        endif
    endwhile
  
```

Inserting a Value in a BST: Pseudo-Code

```

def INSERT_BST( $T, v$ ): # v is the new value
     $x \leftarrow T.\text{root}$ 
     $y \leftarrow \text{NIL}$     # y is x's parent

    # search the right place for z
    while  $x \neq \text{NIL}$ :
         $y \leftarrow x$ 
        if  $v \preceq x.\text{key}$ :
            if  $x.\text{key} \preceq v$ :
                return HANDLE.MULTIINSERT( $x, v$ )
            endif
             $x \leftarrow x.\text{left}$ 
        else:
             $x \leftarrow x.\text{right}$ 
        endif
    endwhile
  
```

$O(h_T)$

Inserting a Value in a BST: Pseudo-Code (Cont'd)

```
# attaching the new node
x ← NEW_NODE(v)
if T.root ≠ NIL:
    if v ≤ y.key:
        SET_CHILD(y, LEFT, x)
    else:
        SET_CHILD(y, RIGHT, x)
    endif
else:
    T.root ← x
endif

return x
enddef
```

Inserting a Value in a BST: Pseudo-Code (Cont'd)

```
# attaching the new node
x ← NEW_NODE(v)
if T.root ≠ NIL:
    if v ≤ y.key:
        SET_CHILD(y, LEFT, x)
    else:
        SET_CHILD(y, RIGHT, x)
    endif
else:
    T.root ← x
endif

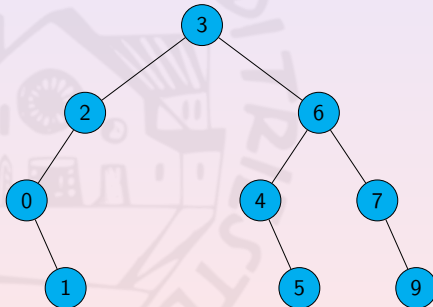
return x
enddef
```

$\Theta(1)$

Removing a Key from a BST

Search the node x containing the key. Either

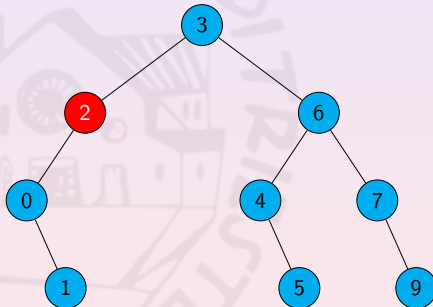
- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either


- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either

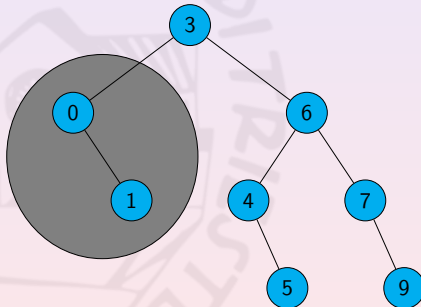
- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either

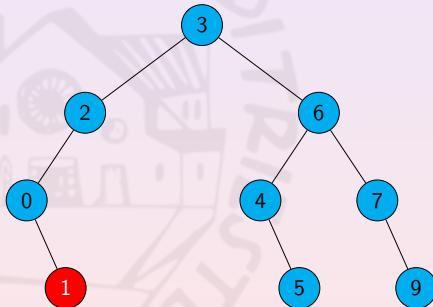
- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either


- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either

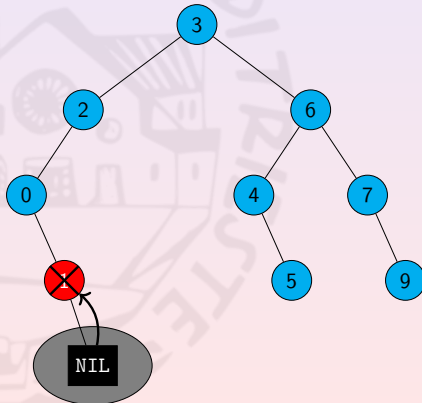
- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either

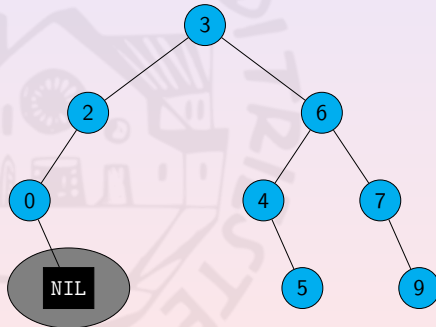
- x has one child at most or



Removing a Key from a BST

Search the node x containing the key. Either

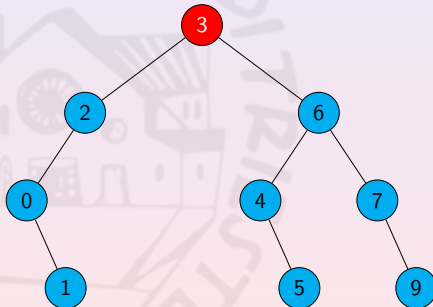
- x has one child at most or



Removing a Key from a BST

Search the node x containing the key. Either


- x has one child at most or
- x has two children



Removing a Key from a BST

Search the node x containing the key. Either


- x has one child at most or
- x has two children



Removing a Key from a BST

Search the node x containing the key. Either


- x has one child at most or
- x has two children



Removing a Key from a BST

Search the node x containing the key. Either

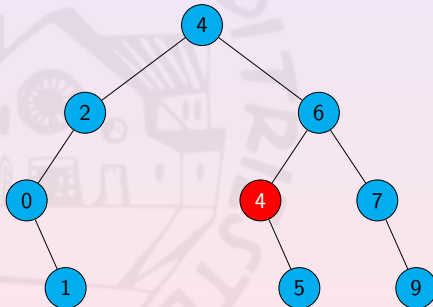
- x has one child at most or
- x has two children




Removing a Key from a BST

Search the node x containing the key. Either

- x has one child at most or
- x has two children

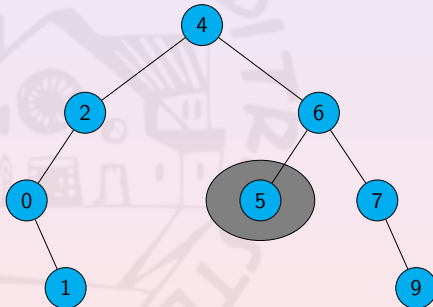




Removing a Key from a BST

Search the node x containing the key. Either

- x has one child at most or
- x has two children



Removing a Key from a BST: Pseudo-Code

```

def TRANSPLANT(T,x,y): # replace x by y
  if IS_ROOT(x):
    T.root  $\leftarrow$  y

    if y  $\neq$  NIL:
      y.parent  $\leftarrow$  NIL
    endif
  else: # x has a parent
    x_side  $\leftarrow$  CHILDHOOD_SIDE(x)

    # attach y in place of x
    SET_CHILD(x.parent, x_side, y)
  endif
enddef

```

Removing a Key from a BST: Pseudo-Code (Cont'd)

```

def REMOVE_BST(T,x): # remove x.key from T and
                     # return a removed node
    if x.left=NIL:    # if x lacks of left child
        TRANSPLANT(T,x,x.right)
        return x
    endif

    if x.right=NIL:   # if x lacks of right child
        TRANSPLANT(T,x,x.left)
        return x
    endif

    y ← MINIMUM_IN_SUBTREE(x.right)
    x.key ← y.key
    return REMOVE_BST(T,y) # y lacks of left child
enddef

```

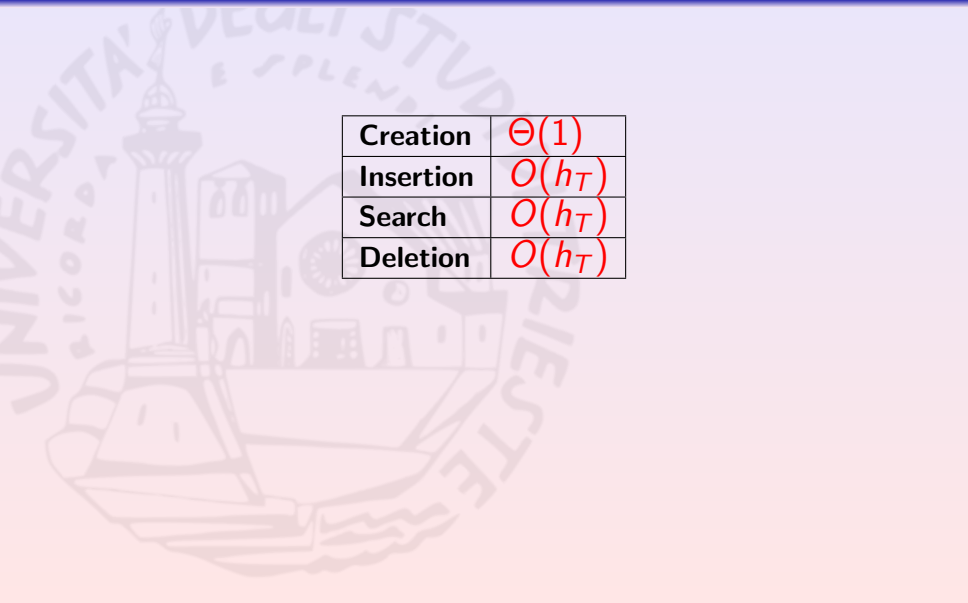
Remove a Node from a BST: Complexity

TRANSPLANT costs $\Theta(1)$, while MINIMUM_IN_TREE $O(h_T)$

So, if x has at most one child, removing it costs $\Theta(1)$

In the general case, removing x takes time $O(h_T)$

Summarizing BSTs...



Creation	$\Theta(1)$
Insertion	$O(h_T)$
Search	$O(h_T)$
Deletion	$O(h_T)$

Summarizing BSTs...

Creation	$\Theta(1)$
Insertion	$O(h_T)$
Search	$O(h_T)$
Deletion	$O(h_T)$

However, h_T may be equal to the number n of nodes e.g., keep inserting always the maximum

Summarizing BSTs...

Creation	$\Theta(1)$
Insertion	$O(n)$
Search	$O(n)$
Deletion	$O(n)$

However, h_T may be equal to the number n of nodes e.g., keep inserting always the maximum

BSTs cost more than single-linked lists (insertion $\Theta(1)$)

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees i.e., bring their heights to $O(\log n)$

How to do it?

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

LOCALLY: balance only the “unbalanced” part of the tree

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

LOCALLY: balance only the “unbalanced” part of the tree
How to know if it is unbalanced?

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance the trees i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

LOCALLY: balance only the “unbalanced” part of the tree
How to know if it is unbalanced? How to handle
branches' lengths?

The background of the slide features a large, faint watermark of the University of Pisa logo. The logo is circular, with the text "UNIVERSITA' DEGLI STUDI DI PISA" around the top and "PISTE" at the bottom. In the center is a detailed illustration of the Leaning Tower of Pisa and other architectural elements.

Red-Black Trees

Are BSTs satisfying the following conditions:

- each node is either a RED or a BLACK node
- the tree's root is BLACK
- all the leaves are BLACK NIL nodes
- all the RED nodes must have BLACK children
- for each node x , all the branches from x contain the same # of black nodes

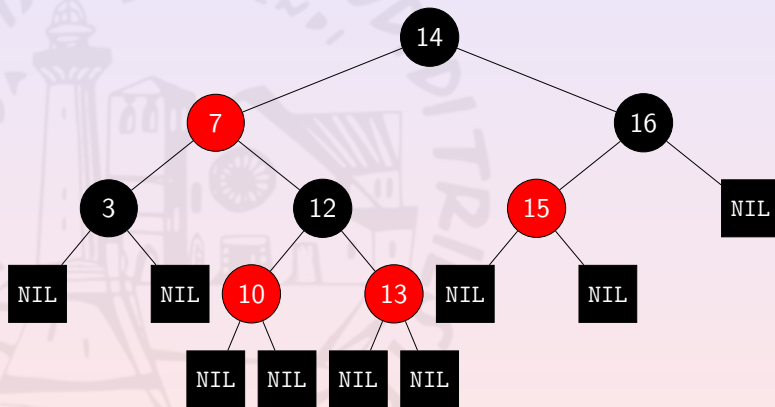
RBTs: Definition

Are BSTs satisfying the following conditions:

- each node is either a RED or a BLACK node
- the tree's root is BLACK
- all the leaves are BLACK NIL nodes
- all the RED nodes must have BLACK children
- for each node x , all the branches from x contain the same # of black nodes

$BH(x)$ will be the # of BLACK nodes below x in any branch

RBTs: An Example



Another Useful $\Theta(1)$ Function

```
def COLOR(x):  
    if x=NIL:  
        return BLACK  
    endif  
  
    return x.color  
enddef
```

How “Tall” Are RB-Trees?

Theorem (Heights of a RB-Tree)

Any RBT with n internal nodes has height at most $2 \log_2(n + 1)$

Proof Sketch:

- prove that the sub-tree rooted in x has at least $2^{BH(x)} - 1$ internal nodes
- $BH(x)$ is at least half of x 's height h then

$$n \geq 2^{h/2} - 1$$

How “Tall” Are RB-Trees?

The ratio between x 's height and $BH(x)$ is topped by 2

Theorem (Heights of a RB-Tree)

Any RBT with n internal nodes has height at most $2 \log_2(n + 1)$


Proof Sketch:

- prove that the sub-tree rooted in x has at least $2^{BH(x)} - 1$ internal nodes
- $BH(x)$ is at least half of x 's height h then

$$n \geq 2^{h/2} - 1$$

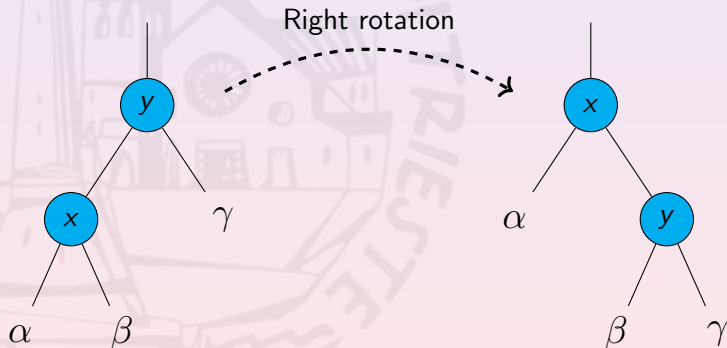
Rotating a Sub-Tree

Rotations are operations on the tree structure. They preserve the BST property.



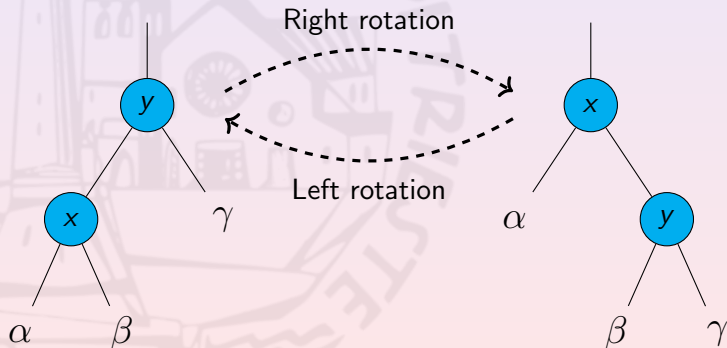
Rotating a Sub-Tree

Rotations are operations on the tree structure. They preserve the BST property.



Rotating a Sub-Tree

Rotations are operations on the tree structure. They preserve the BST property.



Rotating a Sub-Tree: Pseudo-Code

```

def ROTATE(T, x, side):
    r_side ← REVERSE_SIDE(side)

    y ← GET_CHILD(x, r_side)
    TRANSPLANT(T, x, y)

    beta ← GET_CHILD(y, side)
    TRANSPLANT(T, beta, x)

    SET_CHILD(x, r_side, beta) # move beta
enddef
  
```


Rotating a Sub-Tree: Pseudo-Code

```

def ROTATE(T, x, side):
    r_side ← REVERSE_SIDE(side)

    y ← GET_CHILD(x, r_side)
    TRANSPLANT(T, x, y)

    beta ← GET_CHILD(y, side)
    TRANSPLANT(T, beta, x)

    SET_CHILD(x, r_side, beta) # move beta
enddef

```

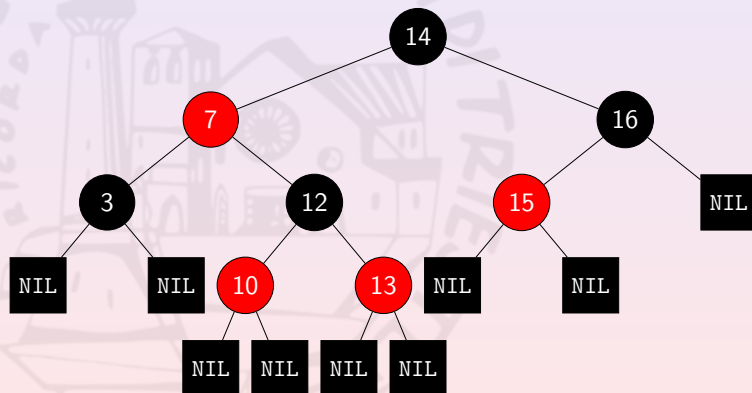
 $\Theta(1)$

Inserting a New Node

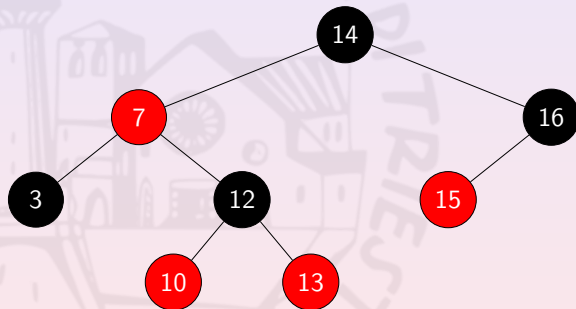
Requires:

- inserting as in BST
- RED-coloring the node
- fixing up RB-Tree properties


Inserting a New Node: Example



Inserting a New Node: Example




Inserting a New Node: Example




Inserting a New Node: Example

x 's parent may be RED. How to fix it?



Inserting a New Node: Example

Case 1: x 's uncle (y) is RED...




```

graph TD
    14((14)) --- 7((7))
    14 --- 16((16))
    7 --- 3((3))
    7 --- 12((12))
    12 --- 10((10))
    12 --- 13((13))
    10 --- 9((9))
    16 --- 15((15))
    style 14 fill:#000,color:#fff
    style 7 fill:#f00,color:#fff
    style 16 fill:#000,color:#fff
    style 3 fill:#000,color:#fff
    style 12 fill:#f00,color:#fff
    style 10 fill:#000,color:#fff
    style 13 fill:#000,color:#fff
    style 9 fill:#f00,color:#fff
    style 15 fill:#f00,color:#fff
    x((x)) --- 9
    y((y)) --- 13
  
```



Inserting a New Node: Example

Case 1: x 's uncle (y) is RED... RED-color x 's granpa and BLACK-color x 's parent and y . New x is former x 's granpa




Inserting a New Node: Example

Still facing problems, but x 's uncle is BLACK (not Case 1)




Inserting a New Node: Example

Case 2: y is BLACK and y and x are on the same side.




Inserting a New Node: Example

Case 2: y is BLACK and y and x are on the same side. Rotate on x 's parent on the opposite side.




Inserting a New Node: Example

Case 2: y is BLACK and y and x are on the same side. Rotate on x 's parent on the opposite side. New x is former x 's parent




Still facing problems, but




Inserting a New Node: Example

Still facing problems, but y is still BLACK (no Case 1)




Inserting a New Node: Example

Still facing problems, but y is still BLACK (no Case 1) and x and y are on different sides (no Case 2)




Inserting a New Node: Example

Case 3: y is BLACK and y and x are on different sides.




Inserting a New Node: Example

Case 3: y is BLACK and y and x are on different sides. Invert x 's pa and granpa colors




Inserting a New Node: Example

Case 3: y is BLACK and y and x are on different sides. Invert x 's pa and granpa colors




Inserting a New Node: Example

Case 3: y is BLACK and y and x are on different sides. Invert x 's pa and granpa colors and rotate on x 's granpa on y 's side



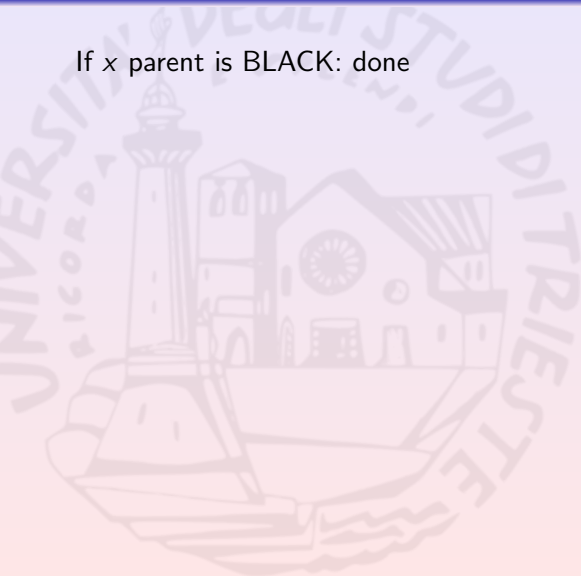
Inserting a New Node: Example

Case 3: y is BLACK and y and x are on different sides. Invert x 's pa and granpa colors and rotate on x 's granpa on y 's side



Inserting a New Node: Complexity

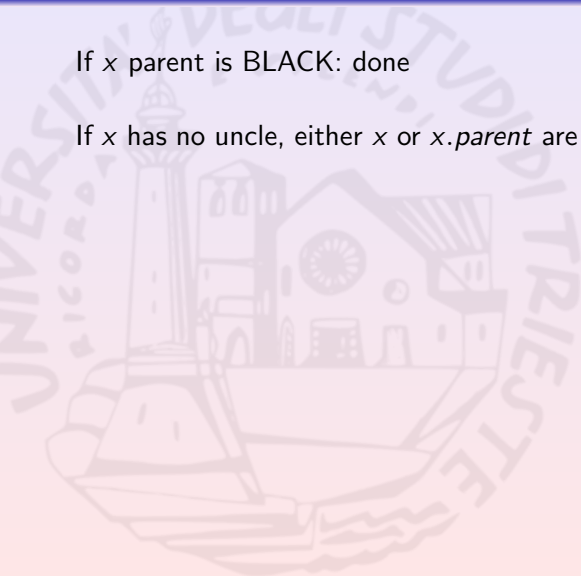
If x parent is BLACK: done



Inserting a New Node: Complexity

If x parent is BLACK: done

If x has no uncle, either x or $x.parent$ are the root: BLACK-color it



Inserting a New Node: Complexity

If x parent is BLACK: done

If x has no uncle, either x or $x.parent$ are the root: BLACK-color it

If x has an uncle, we can always choose between Cases 1, 2, or 3

Inserting a New Node: Complexity

If x parent is BLACK: done

If x has no uncle, either x or $x.parent$ are the root: BLACK-color it

If x has an uncle, we can always choose between Cases 1, 2, or 3

Case 1 either removes the problem or pushes it towards the root

Inserting a New Node: Complexity

If x parent is BLACK: done

If x has no uncle, either x or $x.parent$ are the root: BLACK-color it

If x has an uncle, we can always choose between Cases 1, 2, or 3

Case 1 either removes the problem or pushes it towards the root

Case 2 brings to Case 3

Inserting a New Node: Complexity

If x parent is BLACK: done

If x has no uncle, either x or $x.parent$ are the root: BLACK-color it

If x has an uncle, we can always choose between Cases 1, 2, or 3

Case 1 either removes the problem or pushes it towards the root

Case 2 brings to Case 3

Case 3 solves the problem

Inserting a New Node: Complexity

If x parent is BLACK: done

If x has no uncle, either x or $x.parent$ are the root: BLACK-color it

If x has an uncle, we can always choose between Cases 1, 2, or 3

Case 1 either removes the problem or pushes it towards the root

Case 2 brings to Case 3

Case 3 solves the problem

In the worst case, the algorithm keeps repeating Case 1 steps along the insertion branch and the complexity is $O(\log n)$

Inserting a New Node: Code

```

def INSERT_RBTREE(T, v):
    x ← INSERT_BST(T, v)
    x.color ← RED

    FIX_INSERT_RBTREE(T, x)
enddef

def FIX_INSERT_RBT_CASE1(T, x):
    UNCLE(x).color ← BLACK
    x.parent.color ← BLACK
    GRANDPARENT(x).color ← RED

    return GRANDPARENT(x)
endif

```

Inserting a New Node: Code (Cont'd)

```
def FIX_INSERT_RBT_CASE2(T,x):  
    x_side ← CHILDHOOD_SIDE(x)  
  
    p ← x.parent  
    ROTATE(T,p, REVERSE_SIDE(x_side))  
  
    return p  
endif
```

Inserting a New Node: Code (Cont'd 2)

```
def FIX_INSERT_RBT_CASE3(T, x):  
    g ← GRANDPARENT(z)  
  
    x.parent.color ← BLACK  
    g.color ← RED  
  
    x_side ← CHILDHOOD_SIDE(x)  
  
    ROTATE(T, g, REVERSE_SIDE(x_side))  
endif
```

Inserting a New Node: Code (Cont'd 3)

```

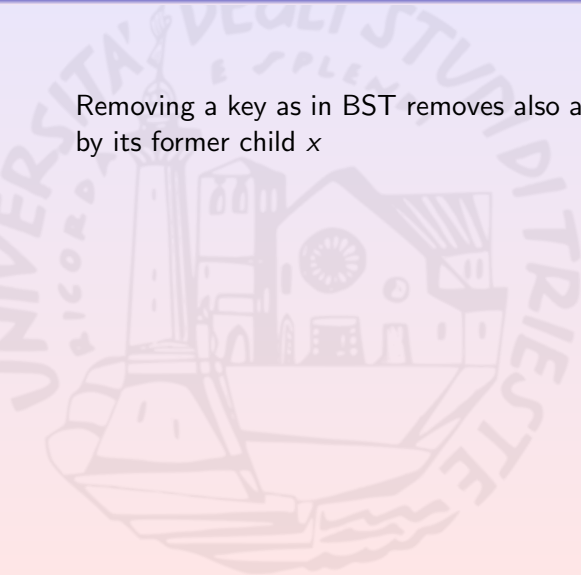
def FIX_INSERT_RBTREE(T,x):
    while (¬IS_ROOT(x) and
           (¬IS_ROOT(x.parent) or x.parent.color=RED)):
        if COLOR(UNCLE(x))=RED:
            z ← FIX_INSERT_RBT_CASE1(T,x)
        else:
            if (CHILDHOOD_SIDE(x) ≠
                CHILDHOOD_SIDE(x.parent)):
                z ← FIX_INSERT_RBT_CASE2(T,x)
            endif
            FIX_INSERT_RBT_CASE3(T,x)
        endif
    endwhile

    T.root.color ← BLACK
enddef

```


Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x



Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x

If y was RED, the RB-Tree properties are preserved

Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x

If y was RED, the RB-Tree properties are preserved

If y was BLACK, the branches through x lost 1 BLACK node

In particular, $BH(x) = BH(w) - 1$ where w is x 's sibling

Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x

If y was RED, the RB-Tree properties are preserved


If y was BLACK, the branches through x lost 1 BLACK node

In particular, $BH(x) = BH(w) - 1$ where w is x 's sibling

The fixing procedure iteratively balances BH on the sub-tree rooted on x 's parent

Removing a Key: Case 0

x is RED




Removing a Key: Case 0

x is RED


- BLACK-color x

$BH(x)$ is increased by 1 and the tree has been fixed



Removing a Key: Case 1

x 's sibling is RED




Removing a Key: Case 1

x 's sibling is RED


- invert colors in x 's parent and sibling
- rotate x 's parent on x 's side

$BH(x)$ does not change



Removing a Key: Case 2

x 's sibling and nephews are BLACK





Removing a Key: Case 2

x 's sibling and nephews are BLACK


- RED-color x 's sibling

$BH(x)$ does not change, while the BLACK height of both x 's parent and sibling are decreased by 1



Removing a Key: Case 3

Among x 's sibling and nephews, only the nephew on x 's side is RED




Removing a Key: Case 3

Among x 's sibling and nephews, only the nephew on x 's side is RED


- rotate x 's sibling on the opposite side w.r.t. x
- invert colors in both old and new siblings of x

The BLACK height of both x and x 's parent does not change



Removing a Key: Case 4

The x 's nephew on the opposite side w.r.t. x is RED




Removing a Key: Case 4

The x 's nephew on the opposite side w.r.t. x is RED

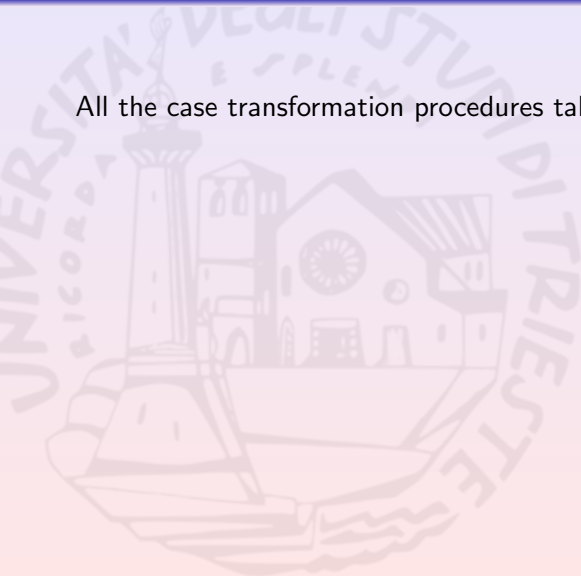
- switch colors of x 's parent and sibling
- BLACK-color the x 's nephew on the opposite side w.r.t. x
- rotate x 's parent on x 's side

The BLACK height of both x and x 's parent does not change



Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$



Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Case 2 pushes the problem one step towards the root

Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Case 2 pushes the problem one step towards the root

If Case 2 occurs after Case 1, Case 0 occurs next

Removing a Key: Some Considerations

All the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Case 2 pushes the problem one step towards the root

If Case 2 occurs after Case 1, Case 0 occurs next

Case 3 transformation procedure brings to Case 4

Removing a Key: Some Considerations

In the worst case scenario, Case 2 is repeated until the problem is pushed up to the root ($O(\log n)$ times)

If this is the case, we have decreased the BLACK height of the tree and the problem is no more a problem

Removing a Key: Code

```

def REMOVE_RBTREE(T, z):
    y ← REMOVE_BST(T, z)

    if y.color = BLACK: # fix from its replacement
        if y.left = NIL
            x ← y.right
        else:
            x ← y.left
        endif
        FIX_REMOVE_RBTREE(T, x)
    endif

    return y
enddef

```

Removing a Key: Code (Cont'd 2)

```
def FIX_REMOVE_RBT_CASE1(T, x):  
    SIBLING(x).color ← BLACK  
    x.parent.color ← RED  
  
    ROTATE(T, x, CHILDHOOD_SIDE(x))  
endif  
  
def FIX_REMOVE_RBT_CASE2(T, x):  
    SIBLING(x).color ← RED  
    return x.parent  
endif
```

Removing a Key: Code (Cont'd 3)

```

def FIX_REMOVE_RBT_CASE3(T,x):
    x_side ← CHILDHOOD_SIDE(x)
    r_side ← REVERSE_SIDE(x_side)

    w ← GET_CHILD(w,r_side)
    GET_CHILD(w,x_side).color ← BLACK
    w.color ← RED

    ROTATE(T,w,r_side)
endif

```


Removing a Key: Code (Cont'd 4)

```

def FIX_REMOVE_RBT_CASE4(T, x):
    x_side ← CHILDHOOD_SIDE(x)
    r_side ← REVERSE_SIDE(x_side)

    w ← GET_CHILD(w, r_side)
    GET_CHILD(w, r_side).color ← BLACK
    w.color ← x.parent.color
    x.parent.color ← BLACK

    ROTATE(T, x.parent, x_side)
endif

```

Removing a Key: Code (Cont'd 5)

```

def FIX_REMOVE_RBT(T, x):
    while x ≠ T.root and x.color ≠ RED:
        w ← SIBLING(x)
        if w = RED:
            x ← FIX_REMOVE_RBT_CASE1(T, x)
        else:
            x_side ← CHILDHOOD_SIDE(x)
            r_side ← REVERSE_SIDE(x_side)

```

Removing a Key: Code (Cont'd 6)

```
    if GET_CHILD(w, r_size) = RED:
        FIX_REMOVE_RBT_CASE4(T, x)

        return
    else:
        if GET_CHILD(w, x_side) = RED:
            FIX_REMOVE_RBT_CASE3(T, x)
        else:
            FIX_REMOVE_RBT_CASE2(T, x)
        endif
    endif
endwhile
enddef
```